

Campus Connect AI

Local Knowledge Chatbot for International Students

Team: Ayana Hussain, Serena Dhillon, Poorvi Bhatia

From Simon Fraser University, BC

Mentors: Lauren Zung & Diane Johnson

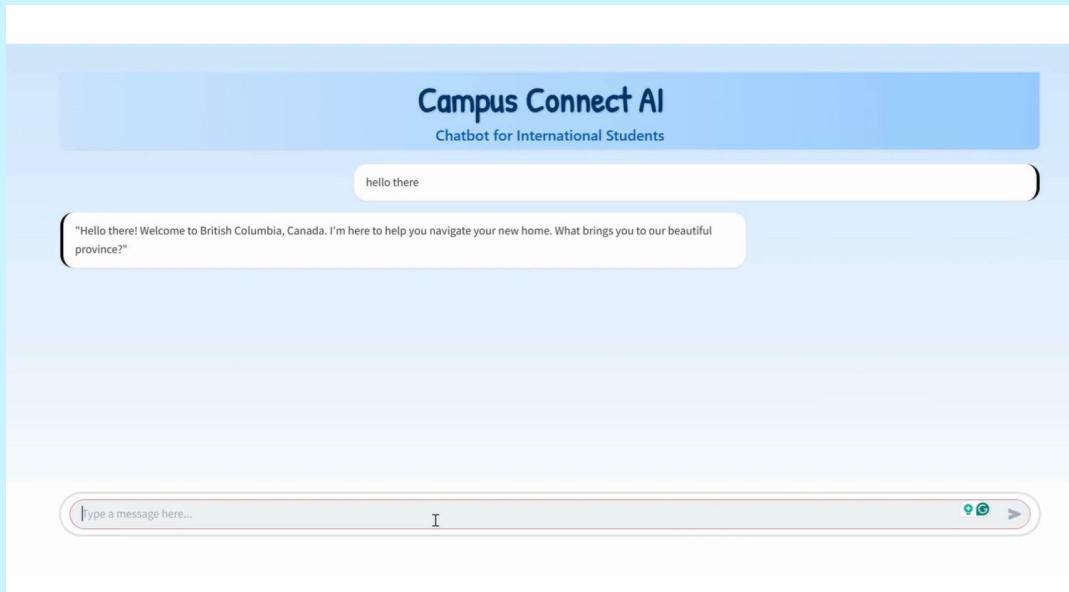


Table of contents

01



Introduction

02



Methods

03



Evaluation

04



Accomplishments +
Future Work

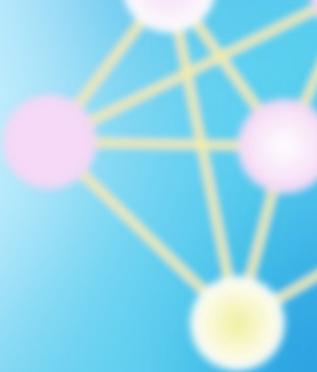
05



Conclusion

01

Introduction



The Problem

Lost in Information: The Challenges International Students Face

- **Key Challenges:** New environments, admin tasks, housing, and healthcare
- **Information Gap**
- **Time Pressure**
- **Not Knowing Who to Ask**



Motivation & Our Vision: Empowering Students Through Intelligent Information Access

Why This Topic Matters

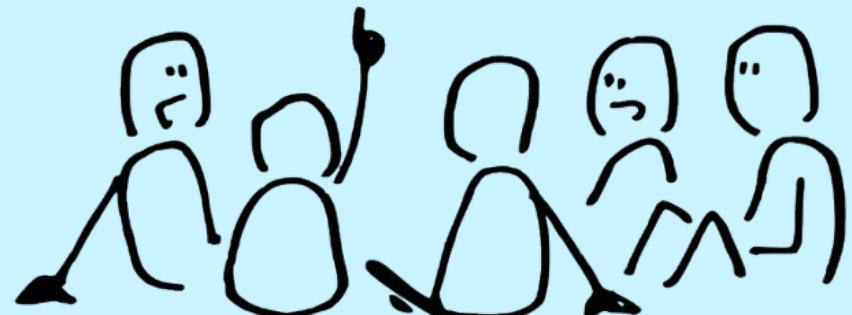
- **Personal experience**
- **Witnessed friends struggle**
- **Observed difficulties accessing timely help when needed**



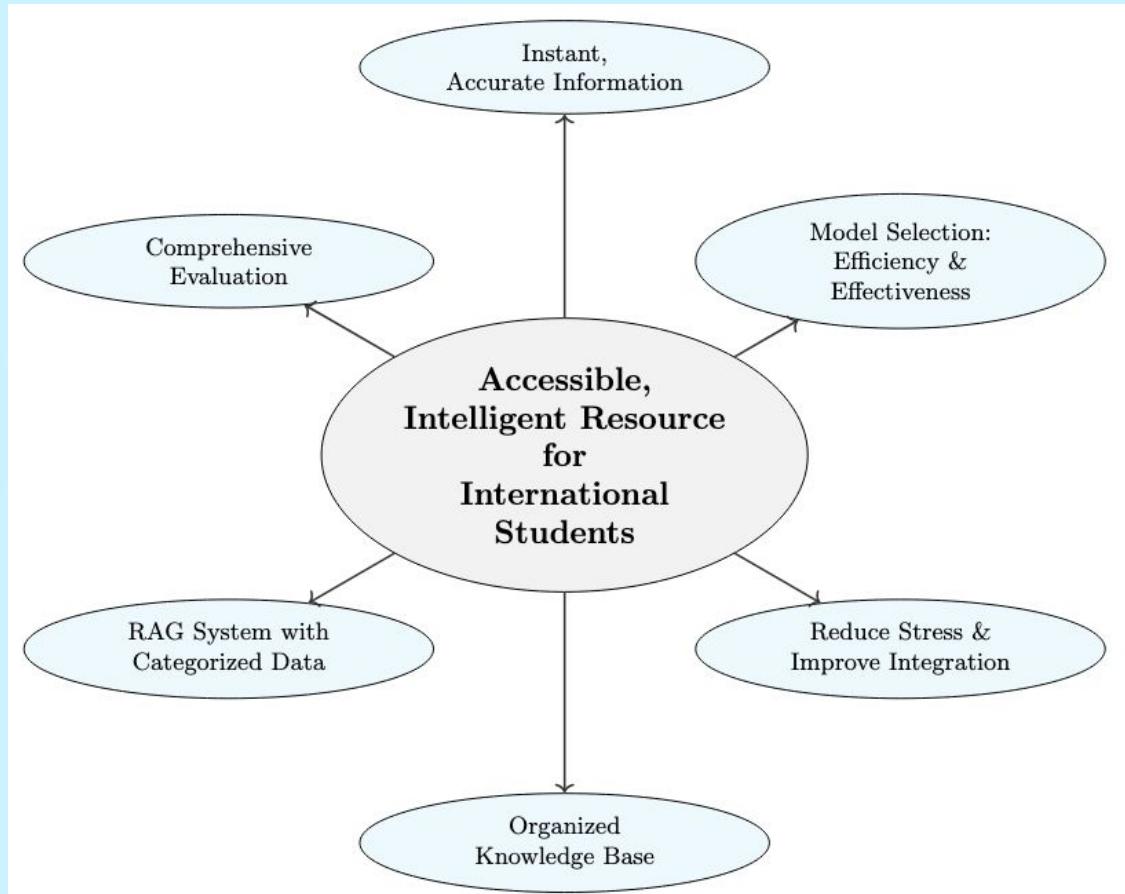
Motivation & Our Vision: Empowering Students Through Intelligent Information Access (Cont)

Benefits of Our Chatbot

- Provides **24/7 assistance**
- Simplifies **access to critical information**
- Allows students to **focus on academic success**
- **Inclusive campus experience**



Goals



02

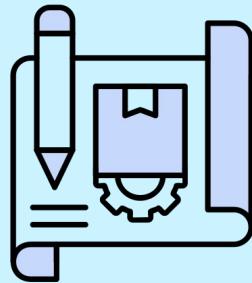
Methods



Building the System: Our Iterative Approach

01

Initial Prototyping



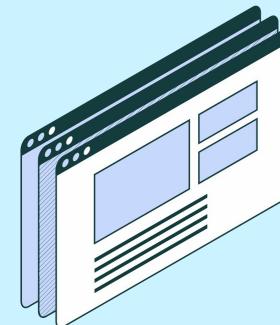
02

Model
Experimentation



03

User Interface



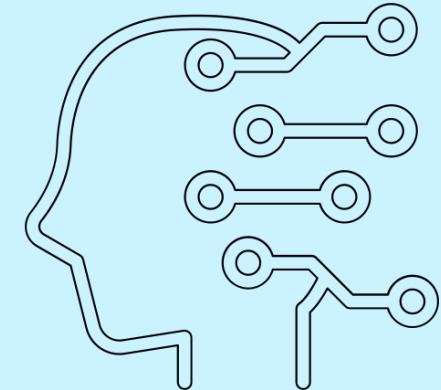
Data Domains

- **University Specific**
- **Immigration**
- **Healthcare**
- **Transit**
- **Activities**
- **Cost of Living**



Models & Development

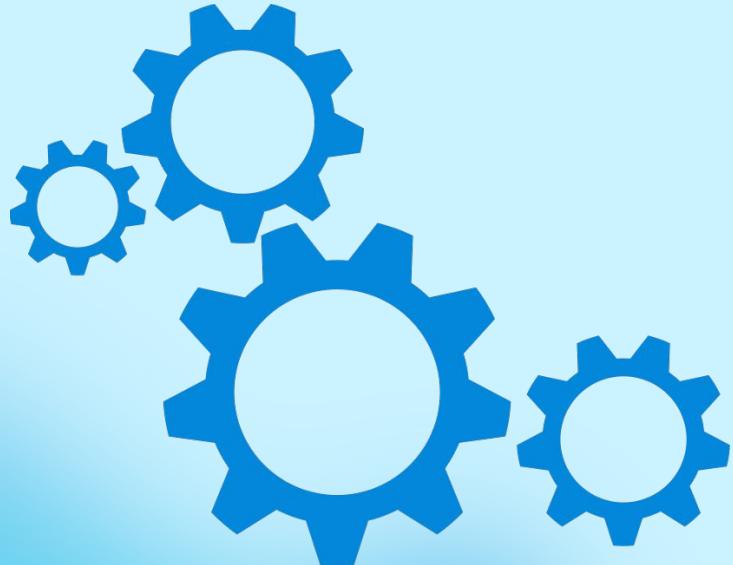
| Embedding Model | Time Taken |
|--------------------------------------------------|--------------|
| sentence-transformers/paraphrase-MiniLM-L6-v2 | 30.5s |
| sentence-transformers/all-MiniLM-L6-v2 | 32.7s |
| sentence-transformers/multi-qa-MiniLM-L6-cos-v1 | 49.5s |
| sentence-transformers/all-MiniLM-L12-v2 | 54.4s |
| sentence-transformers/all-roberta-large-v1 | 1m 55.8s |
| sentence-transformers/multi-qa-mpnet-base-dot-v1 | 2m 18.7s |
| neuml/pubmedbert-base-embeddings | 2m 27.4s |
| sentence-transformers/paraphrase-mpnet-base-v2 | over 3m |



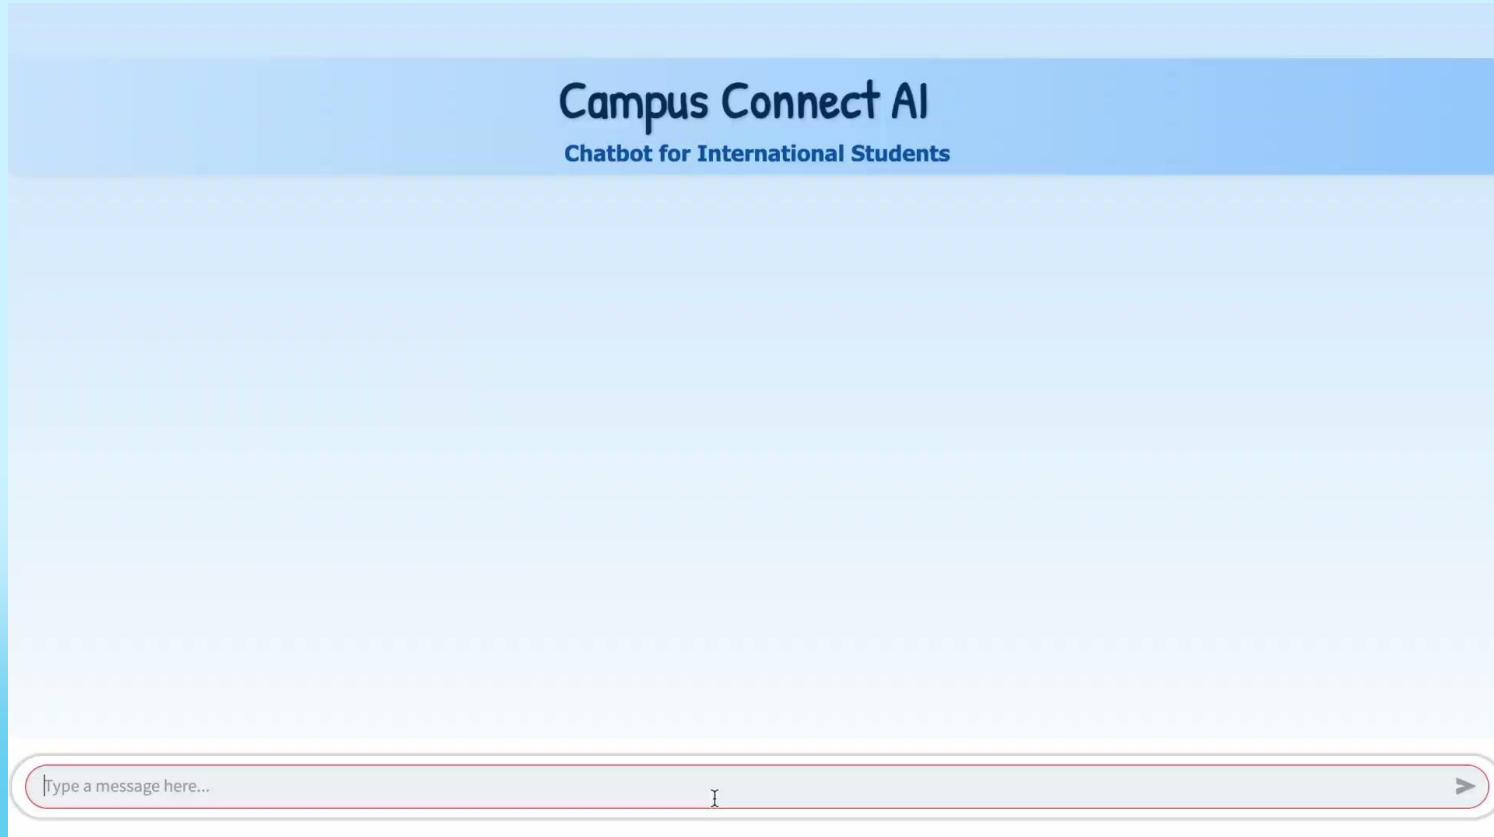
Instruction-Tuned Model Tested

google/gemma-2-2b-it
mistral-7b-instruct-v0.2.Q4_0.gguf
meta-llama/Llama-3.2-1B-Instruct
.Llama-3.2-1B-Instruct-IQ3_M.gguf

Prompt Engineering



Interface Demo



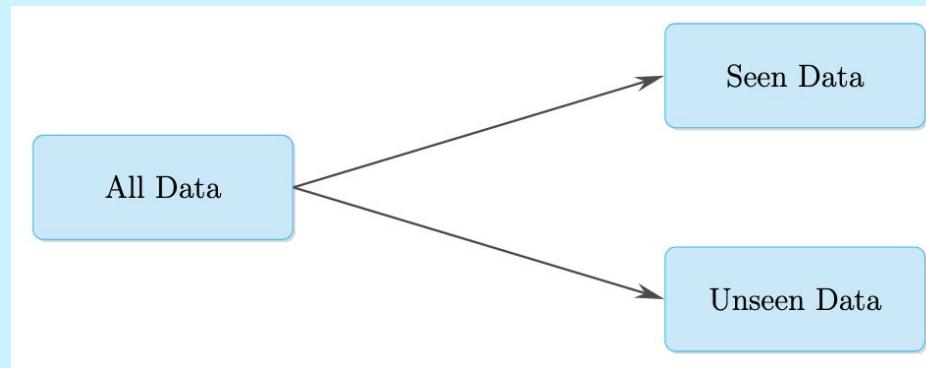
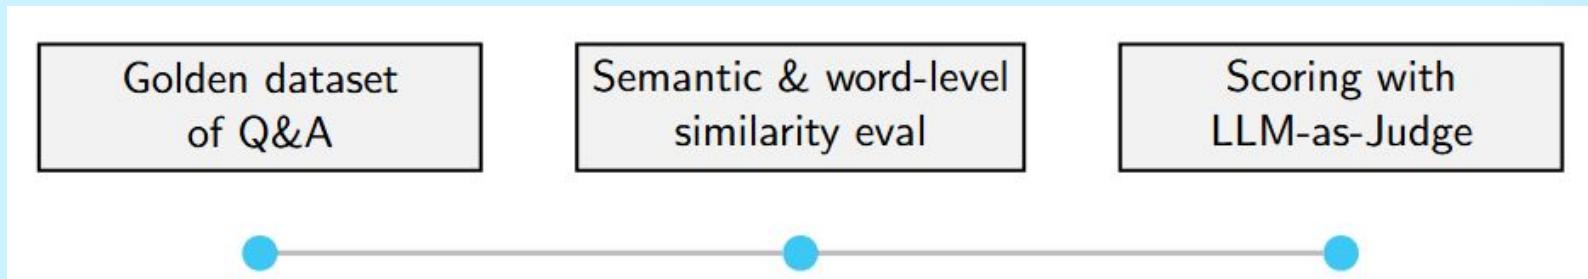
03

Evaluation



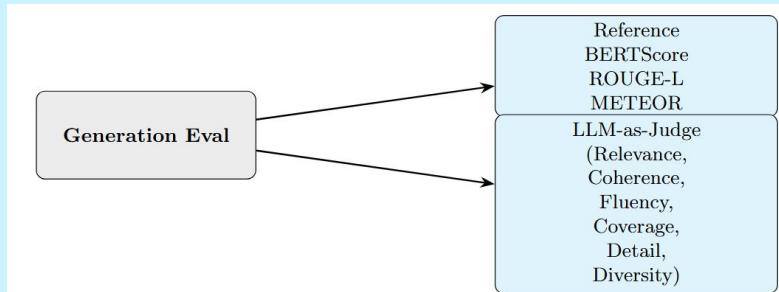
Evaluation Approaches

Automated Evaluation

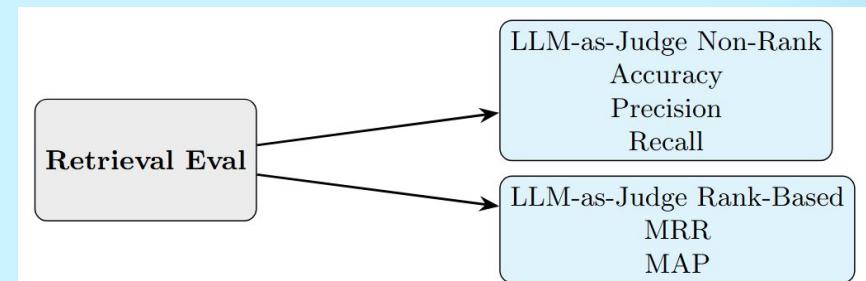


Evaluation Metrics

Generation Model Quality



Retrieval Model Quality



Generation Evaluation Results + Takeaways

Seen Data

- LLM Evaluation

| Criterion | Score |
|-----------------|-------|
| Relevance | 4.92 |
| Coherence | 4.81 |
| Fluency | 5.00 |
| Coverage | 3.96 |
| Level of Detail | 2.92 |
| Diversity | 1.19 |

- Comparison Metrics

| Metric | Score |
|---------|--------|
| BERT F1 | 0.9057 |
| METEOR | 0.4170 |
| ROUGE-L | 0.3748 |

- Takeaway

- High quality answers and contextual embedding similarity despite different wording than the reference answers.

Generation Evaluation Results + Takeaways

Unseen Data

- LLM Evaluation

| Criterion | Score |
|-----------------|-------|
| Relevance | 4.74 |
| Coherence | 4.65 |
| Fluency | 4.96 |
| Coverage | 3.91 |
| Level of Detail | 2.91 |
| Diversity | 1.26 |

- Comparison Metrics

| Metric | Score |
|---------|--------|
| BERT F1 | 0.8925 |
| METEOR | 0.3318 |
| ROUGE-L | 0.2750 |

- **Takeaway**

- Overall responses scored as very high quality (with a small drop compared to seen data)

Retrieval Evaluation Results + Takeaways

Seen Data

| Metric | Score |
|-----------|-------|
| Accuracy | 0.825 |
| Precision | 0.804 |
| Recall | 0.686 |
| MRR | 0.941 |
| MAP | 0.941 |

Unseen Data

| Metric | Score |
|-----------|-------|
| Accuracy | 0.319 |
| Precision | 0.319 |
| Recall | 0.194 |
| MRR | 0.417 |
| MAP | 0.361 |

- **Takeaway**
 - Seen data: LLM judges documents as relevant most of the time
 - Unseen data: Potential to add more documents
 - LLM is able to use relevant info and ignore irrelevant content in the documents

04



Accomplishments + Future Work



Accomplishments

- Built **end-to-end RAG prototype**
- **Cleaned and integrated data**
- Developed a **Streamlit UI**
- Achieved **high generation quality scores**
- For seen data, **documents retrieved** are both **accurate** and **well-ordered**
- For unseen data, we had **high quality responses** although there was **room for improvement**



Learning

RAG Systems and Implementation

- Practical experience with Hugging Face, Streamlit, and Ollama
- Impact of embedding models on performance and retrieval
- Model quantization techniques
- Prompt engineering techniques

RAG Evaluation

- RAG Evaluation
- Gold Benchmark Datasets



Challenges

Data Processing

- Optimal data chunk size for retrieval
- Selecting relevant data sources
- Data cleaning, aggregation, and formatting

Model Performance

- Embedding models
- Embedding quality
- Overall system runtime
- Resource limitations

System Implementation

- UI
- Storing the Documents

Future Directions

01 —————• 02 —————• 03 —————•

User Testing



Add more data

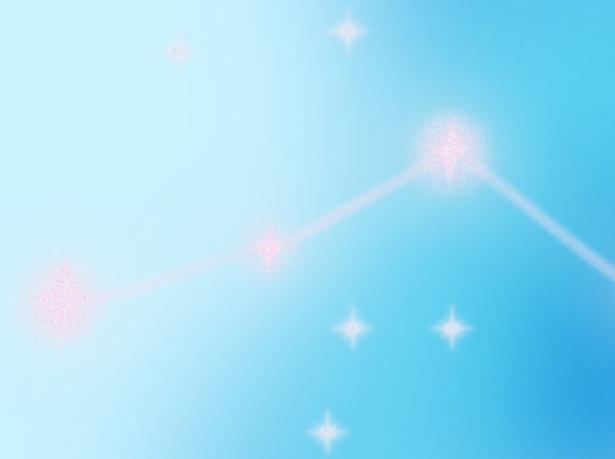
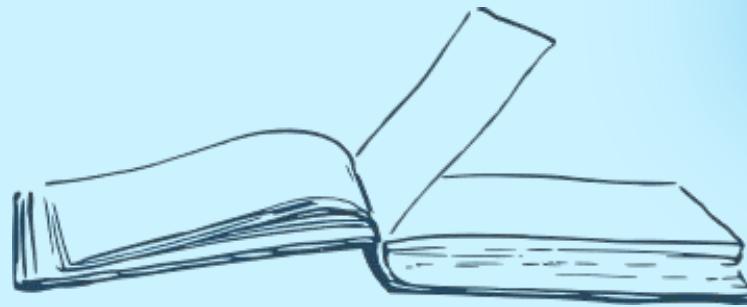


**Deploy to
Hugging Face**



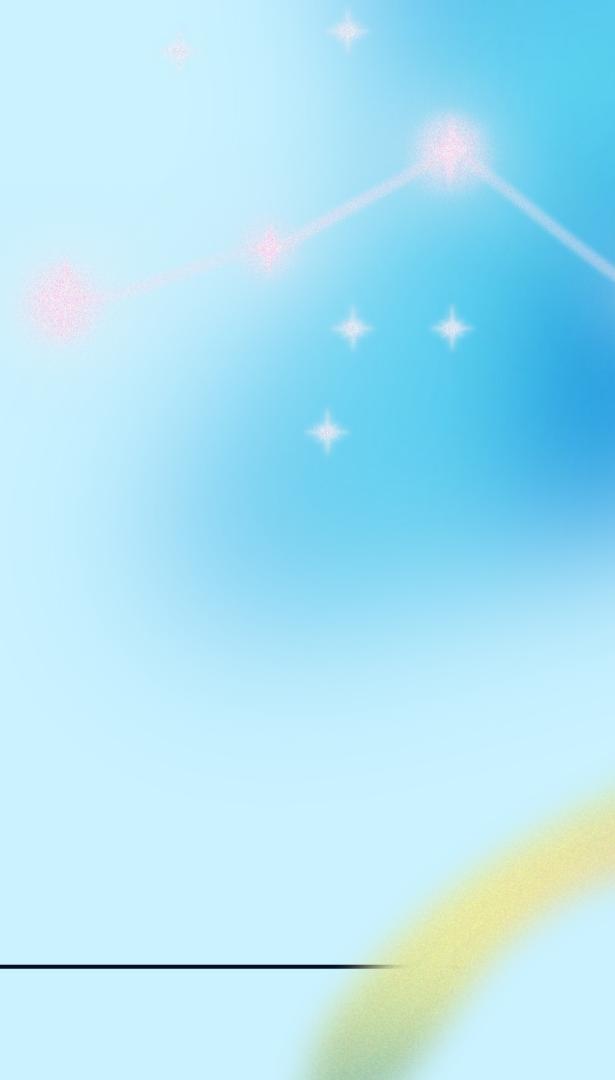
05

THE END



06

Appendix



Prompt Engineering

Initial Prompt

```
rag_prompt = f""""
```

```
You are a helpful assistant for international students new to British Columbia Canada. Here are relevant documents:
```

```
{relevant_texts}
```

```
Please respond to the following question. Be conversational but concise, aim to answer accurately using the documents,  
but in as few words as possible (i.e. less than 20).
```

```
DO NOT USE THE DOCUMENTS IF THEY ARE NOT HELPFUL FOR THE QUERY.
```

```
Do not ask the user irrelevant questions unless it relates to their query.
```

```
Question: {query}
```

```
Answer:
```

```
""""
```

Category-Wise Prompts

```
# category-wise prompts
hike_prompt = """
INSTRUCTIONS:
  1. Convert structured information about the hike into a short, friendly paragraph using natural language.
  Do not repeat numbers or use formatting from the source.
  2. If I may ask about hiking information, only answer with required information. Users can ask for more information if needed.
  3. When asked for a particular type of hike, find it instead of saying that one would not work in the category they asked for.
  4. DO NOT list trail attributes or stats (like "Distance: 3.1 km, Elevation: 789 m").
  Instead, describe them in context (e.g., "a steep 3 km trail with a tough 789 m climb").
  5. Avoid repeating exact numbers unless essential (e.g., elevation gain is helpful, but don't dump all stats).
```

food_prompt = """
INSTANTANEOUS
1. Convert structured food and dining information into a friendly, helpful paragraph. Do not copy the question or use list formatting.
2. Only answer what the user asked. Do NOT add anything that wasn't requested.
3. Describe details in a natural way (e.g., "open 24/7" during the semester) instead of "Hours: 24/7").
4. Mention unique features only when they help clarify the user's question.
5. If the question asks for a specific place (e.g., a meal plan, or food station), describe it clearly in context.
6. If the question can't be answered from the data, respond with:
 - "I'm sorry, I don't have that information. Please check the official SFU Food website."
7. Provide the official link when available and relevant to the answer.
8. Do **Not** list menu items, prices, or square footage unless directly relevant to the user's question.
9. Only provide food information that is relevant. If they ask for some place that serves a chicken sandwich do not provide information to a vegan place.

```
housing_prompt = """  
INSTRUCTIONS:  
1. Provide structured information about SFU or general student housing into a short, friendly paragraph using natural language.  
Do not repeat, formatting or list prices unless helpful for context.  
2. Focus on what matters to the student: location, room types, meal plans, how to apply, and support available.  
3. Only mention costs in a general way (e.g., "starts around $4,000 per term") unless the user explicitly asks for detailed pricing.  
4. If information varies (e.g., by room type or campus), explain this clearly but briefly.  
5. If the user asks a specific housing question and the answer depends on certain conditions  
(e.g., term length, student status), explain those conditions clearly and simply.  
6. If the answer is not known or not in the data, respond with:  
| "I'm sorry, I don't have that information. Please check the SFU Housing website for details."  
7. DO NOT DUMP full lists of buildings, prices, or amenities. Summarize and keep it conversational.  
8. If the information is specific to SFU, make sure you say it to be clear.
```

```
parks_prompt = """  
INSTRUCTIONS:  
    1. Convert structured information about the park into a short, friendly paragraph using natural language.  
        Do not repeat numbers or use formatting from the source.  
    2. Provide only necessary information that will allow the user to enjoy the park.  
        - Feel free to tell them about logistical information if asked.
```

```
# activities general: covers how to answer general parks, hikes, food, clubs, cultural related questions
activities_general = f"""
INSTRUCTIONS:
    1. If they ask for suggestions, provide 2 to 3 suggestions.
    2. Do NOT list all information. Instead describe them in context
    3. Provide accurate suggestions, NOT suggestions of things that will not work for what they want.
    4. Convert structured information about the activity into a short, friendly paragraph using natural language.
    Do not repeat formatting from the source.
```

```
# permits prompt: covers ways to answer immigration, study permits, work permits, and permanent residence related questions
permits_prompt = """  
INSTRUCTIONS:  
    1. When given a specific question with many possible answers, you can ask for more specific information.  
        | - if they are not asking for an extension do not provide information in regards to an extension of a permit.  
    2. Only answer with information provided  
        | - Information should NOT be guessed and DO NOT add extra information  
    3. If the answer is not in the dataset, respond with: "I'm sorry, I don't have that information.  
        Please check the official IRCC website for more details."  
    4. If it is helpful, provide the link and a description about it.  
    5. Do NOT list all information. Instead describe them in context  
    6. If the answer depends on a specific condition explain those clearly.  
    7. Do NOT make assumptions about the user's situation.
```

```
transit_prompt = """  
INSTRUCTIONS:  
1. Convert structured information about public transit into a short, friendly paragraph using natural language.  
2. Do not list statistics or technical formatting (like route numbers or fare counts) unless directly relevant to the user's question.  
3. Encourage transit options clearly – describe them in context (e.g., "a quick SkyTrain ride from downtown to the airport").  
4. Provide only what the user needs to understand how to get around or plan their trip.  
5. If the user is asking for directions, give a general summary of how they might travel.  
6. If the question is about fares, schedules, or route planning and the exact info is not available,  
| tell the user to check the TransLink website and briefly explain what they can find there.  
7. DO NOT guess or make up transit information.  
8. If the information is not in the source, say "I'm sorry, I don't have that information. You can check the official TransLink site for more details."  
"""
```

Main Final Prompt

```
# main rag prompt
rag_prompt = f"""
You are a helpful, friendly assistant for international students new to British Columbia, Canada.
```

Below are some reference documents that may be relevant to the user's question:
{relevant_texts}

INSTRUCTIONS:

1. If the user's query is just a greeting (like "hello", "hi", "what's up"):

- Respond with a single brief friendly greeting
- Offer to help with questions about studying or living in BC
- Do NOT include ANY information from the reference documents
- Do NOT create additional answers beyond answering their original question

2. If the user is asking for information:

- Be friendly and answer based ONLY on the reference documents if relevant
- Summarize the necessary information into a couple sentences.
- Do NOT create additional questions and answers beyond answering their original question
- Limit your entire response to no more than 3 concise sentences when possible. Do not create long multi-line answers.
- If the documents don't provide sufficient information, say "I don't have enough information to answer that. Please refer to official sources."
- Ask for more information when there are multiple scenarios in the documents.
- If they ask things like "can I", "will I", "how can I" feel free to ask follow up questions if you don't know how to answer with the information provided. Do not just assume.

3. IMPORTANT: Never generate additional content beyond answering the user's question. Do NOT number or bullet your points. Always use natural sentences and group similar information together where possible.

User question: {query}

Your response (just the answer, no preamble):

"""

Credits

Title slide icons:

<https://www.slidescarnival.com/template/linguistic-philosophy-of-education-slides/202625>

Blue and pink graphics:

<https://www.slidescarnival.com/template/book-club-marketing-plan/29147>

Music track: Tech by Aylex

Source: <https://freetouse.com/music>

No Copyright Background Music